

Anthony Ge

+1 (914) 770-5675 | anthonykounge@gmail.com | github.com/SomeGuy966 | linkedin.com/in/anthony-ge/

Education

Northwestern University | Evanston, IL

Expected June 2028

Bachelor of Science, Computer Science and Mathematics (Double Major)

GPA: 3.95/4.00

- Courses: Probability & Stochastic Proc, ML, Applied Statistical Learning, Data Structures & Algorithms, Real Analysis, Optimization

Experience

Amazon

Jun. 2026 – Sep. 2026

SDE Intern, Amazon Prime Video

Seattle, WA

- Cut defect-discovery latency on Prime Video's catalog export feeds from days or weeks to ~5 minutes as sole author of an event-driven data-integrity monitoring service (Java 17, AWS Lambda, SNS→SQS; 12 CloudWatch metrics, 14 auto-ticketing alarms)
- Surfaced 2,465 XSD schema violations in the first 3 hours of production traffic and 17,152 missing-required-field violations traced to a producer-side gap; verified against production bytes that 98.3% of 4,848 violations sat in a single field
- Delivered ~16,600 lines across two services (56% tests; 398 tests behind a 98% line/branch coverage gate) in 51 solo code reviews over ~7 weeks; all infrastructure in TypeScript CDK across 4 accounts and 2 regions on one self-mutating pipeline

MedSphere

Jan. 2026 – Mar. 2026

Software Engineer Intern, AI

Skokie, IL

- Built document-to-graph pipeline extracting pharma/biotech text/metadata into graph nodes and edges
- Implemented entity extraction and normalization (canonicalization, deduping, synonym resolution) for consistent schema at scale
- Built relation extraction and graph construction for confidence-weighted typed edges and source-backed explanations
- Productionized graph builds via async jobs for batch/incremental updates, plus logging/monitoring and stable APIs for consumers

Projects/Competitions

BTC-USD Order-Flow Research Pipeline (C++/Python)

github.com/SomeGuy966/btc-trading-pipeline

- Built a hybrid C++/Python research pipeline for Gemini BTC-USD: REST ingest (Python and C++ implementations behind one pybind11 interface), header-only C++ order-flow features, and online Lasso inference refit every tick on a rolling 10-tick window
- Profiled the Python/C++ boundary: per-feature pybind11 calls ran at $0.66\times$ pure Python; batching all four features into a single crossing and passing trades by const reference reached $2.32\times$ Python ($3.5\times$ over per-call), equivalence-verified before timing
- Diagnosed and fixed a state-leakage bug where a rolling-window feature was re-fed on every model refit; added regression tests
- Built JSONL record/replay for deterministic offline runs and a one-command CMake/Conan/Ninja/Poetry build; four GitHub Actions gate every PR on pytest, ruff/mypy strict, GoogleTest, and clang-format/clang-tidy (85 tests)

Kalshi Weather-Market Forecasting | Python, LightGBM, pytest, mypy-strict, CI

github.com/SomeGuy966/weather-bot

- Built a leakage-safe pipeline turning hourly METARs and NOAA NBM/GFS guidance into calibrated daily-high distributions for 8 U.S. cities with Kalshi markets: 586K hourly snapshots (2020–26) via strict as-of joins and chronological train/calibrate/test splits
- Trained a calibrated 31-class LightGBM model of the high vs. a forecast/running-max anchor, cutting log-loss 30% against NBM's own probabilistic guidance (1.49 vs. 2.12; ECE 0.004); mapped it onto Kalshi contract ladders for real-time partial-day inference
- Archived 3.2K contracts (order-book candles, trade tape) and backtested 52K contract-hours after fees; the market was sharper at every hour and city (Brier 0.120 vs. 0.141); shipped as a documented negative result: paper-trading ledger, no capital deployed

S&P 500 ARIMA vs. Random Walk – Walk-Forward Forecast Evaluation

github.com/SomeGuy966/sp500_arima

- Built a walk-forward evaluation of ARIMA forecasts on 50 years of daily S&P 500 data (600 monthly re-fits, 12.6k out-of-sample days), benchmarked against random-walk and random-walk-with-drift baselines using Diebold–Mariano tests
- Showed the AIC-selected ARIMA's 2.48% monthly MAPE is statistically indistinguishable from a random walk (DM $p = 0.93$) and significantly worse at a one-day horizon ($p = 0.001$); 95% prediction intervals cover only 91% of one-day outcomes
- Backtested a daily sign-of-forecast strategy with transaction costs (gross Sharpe 0.70 vs. 0.60 buy-and-hold, falling to 0.21 at 5 bps) and traced the gross edge to the index's lag-1 autocorrelation flipping from +0.08 (1975–89) to –0.12 (2005–24)
- Engineered as a typed Python package (statsmodels, pandas): CLI, parallel walk-forward engine, 59 tests at 95% coverage, mypy --strict, GitHub Actions CI on Python 3.11–3.13

VIX Probabilistic Forecasting – Walk-Forward Evaluation

github.com/SomeGuy966/IMPLIED-volatility-forecasting

- Built a Python package (arch) forecasting the CBOE VIX 1 and 5 days ahead as full predictive distributions: random walk, EWMA, HAR, AR(1)-GARCH(1,1)-t and HAR-GARCH(1,1)-t, each refit at 676 no-look-ahead walk-forward origins (2024–26)
- Evaluated with Diebold–Mariano (HAC/HLN), Kupiec and Christoffersen coverage tests, Winkler score, and pinball loss
- Found RMSE gains over the random walk of only 2.5% at $h=1$ (HAR-GARCH-t; significant under absolute loss only) and 8.5% at $h=5$ (HAR, DM $p < 0.05$); only the GARCH-t intervals pass Christoffersen independence, scoring 6–10% better
- Shipped as a typed package: CLI, joblib-parallel backtest (3,380 fits/horizon in ~20 s), 62 pytest tests incl. a no-look-ahead test per model and a Monte Carlo size check of the DM test, mypy --strict, CI, byte-for-byte reproducible results

Skills

Languages: Python, C/C++, Java, MATLAB, TypeScript/JavaScript, SQL, Lua

Web & Back-End: Node.js, React.js, Express.js, Svelte, REST APIs, HTML/CSS

Tools/Libraries: PyTorch, TensorFlow, LightGBM, CatBoost, statsmodels, pandas, NumPy, pybind11, Git, GitHub Actions, AWS, Firebase